



Livio AI Dashboard

Usage Analytics & Guardian System

Internal Overview Document

Prepared for CEO & HR Review

Prepared by: Ashwin Pawar, Chief AI Officer

Date: May 5, 2026

Confidential — For Livio Internal Use Only

■ Executive Summary

Livio has deployed a comprehensive AI usage monitoring and protection system. This document outlines three key components: a **real-time Dashboard** for employee usage analytics, a **Watchdog agent** that detects 14 types of anomalies, and a **Fixer agent** that automatically resolves issues — all at **zero additional token cost**.

35	14	\$0	5 min
Employees Monitored	Anomaly Types	Token Cost	Scan Interval

■ 1. OpenClaw Dashboard

A web-based analytics dashboard deployed on GitHub Pages that provides real-time visibility into AI usage across all 35 Livio employees.

Three Views:

- ■ **My Dashboard** — Personal stats — each employee sees their own token usage, weekly hours, badge progress, and rank.
- ■ **HR Dashboard** — Password-protected view for HR — team-wide analytics, bar charts, employee breakdown table, cost tracking.
- ■ **Leaderboard** — Gamified rankings — podium for top 3, weekly vs all-time views, power user badges.

Key Features:

- Real-time data — auto-refreshes every 60 seconds
- Gamification — badges (Beginner → Champion), streaks, rocket animation for 10hr goal
- Cost transparency — per-employee and company-wide token spend in dollars
- Mobile responsive — works on phone, tablet, and desktop
- Zero maintenance — static site, no server costs

■ URL: <https://sweden-technologies.github.io/openclaw-dashboard/>

■ ■ 2. Watchdog Agent

An automated monitoring system that scans all 39 OpenClaw workspaces every 5 minutes, detecting 14 types of anomalies across cost, behavior, and system health.

14 Monitored Metrics:

#	Metric	Threshold	Severity
1	Workspace cost in 1 hour	> \$2.00	■ Critical
2	Single session tokens	> 500,000	■ Critical
3	Daily company spend	> \$10.00	■ Warning
4	Single session cost	> \$0.50	■ Warning
5	Daily per-employee budget	Configurable	■ Warning
6	Same tool called 10+ times (loop)	10x same args	■ Critical
7	Consecutive tool failures	> 5	■ Warning
8	Idle session (no user message)	> 30 min	■ Warning
9	Zombie session (>2hr, <5 interactions)	2hr / 5 int	■ Warning
10	Heavy usage at unusual hours	11pm-7am IST	■ ■ Info
11	User switches to expensive model	Repeated	■ Warning
12	VPS RAM usage	> 85%	■ Critical
13	VPS CPU sustained	> 90% for 10min	■ Warning
14	Container restarts	> 3 in 1 hour	■ Critical

■ 3. Fixer Agent

When the Watchdog detects an issue, the Fixer agent automatically steps in to resolve it. No human intervention needed — the entire pipeline is autonomous.

How It Works:

- 1■■ Watchdog scans — every 5 minutes, monitors all 14 metrics across 39 workspaces
- 2■■ Issues detected — anomalies are logged with severity (critical/warning/info)
- 3■■ Fixer triggered — automatically applies the appropriate fix based on issue type
- 4■■ Fix applied — session compacted, paused, killed, or cleaned up
- 5■■ Re-scan — monitor runs again to verify the fix worked
- 6■■ Dashboard updated — status changes from ■■ to ■ in real-time

Auto-Fix Actions:

Issue Detected	Fixer Action	Approval Needed?
Session file > 10MB	Compact to last 50 messages	No — automatic
AGENTS.md > 20KB	Trim to 6KB	No — automatic
Old checkpoint files	Delete old checkpoints	No — automatic
Cost spike > \$2/hr	Pause workspace sessions	No — automatic
Tool loop detected	Kill looping session	No — automatic
Idle session > 30min	Mark as inactive	No — automatic
RAM > 85%	Clear caches + cleanup	No — automatic

■ 4. Cost & Efficiency

The entire Guardian system operates at **zero token cost**. All monitoring and fixing is done with pure Python scripts — no AI model calls are made by the Watchdog or Fixer.

Component	Technology	Token Cost	Server Load
Dashboard	Next.js (static)	\$0	None — GitHub Pages
Watchdog Monitor	Python script	\$0	~2s every 5 min
Fixer Agent	Python script	\$0	~1s on demand
Data Generator	Node.js script	\$0	~1s every 5 min
Total Additional Cost	—	\$0/mo	Negligible

The only costs are from employees using their OpenClaw agents for daily work. The Guardian system actually **saves money** by catching runaway sessions, infinite loops, and cost spikes before they escalate.

■■ 5. Architecture Overview

Data Flow:

- OpenClaw Sessions (39 workspaces)
↓ every 5 minutes
- Watchdog Monitor (Python) → detects anomalies
↓ if issues found
- Fixer Agent (Python) → applies fixes → re-scans
↓ always
- Data Generator (Node.js) → usage.json + guardian.json
↓ pushed to
- GitHub Pages Dashboard → auto-refreshes every 10s